


Lecture 18 DNS



Administration

- PA7 due tomorrow
- **FINAL PROJECT** writeups due Wednesday
- Midterm #2 a week from today (March 29)
 - You may bring one *double-sided* 8.5x11 set of handwritten notes
 - Content is anything in lecture since Feb 8
 - We'll have review in the 2nd part of lecture on Wednesday; bring questions!

2

The Web at Scale

- In the first third of this class we focused on point-to-point basics of the Web
 - HTML/HTTP basics
 - Sessions & server architecture
 - Crypto and XML
- Middle third: managing lots of Web activity, with emphasis on *semantics*
 - Web search and page management
 - Recommenders, auctions, data mining
- Next, building Web-scale *systems*

3

A Mini Syllabus

- Web is only interesting at massive scale
 - What's data mining without the data?
- Upcoming topics:
 - DNS (today)
 - Scaling and Distribution
 - Caching & Proxies
 - Google File System, MapReduce
 - Cloud Computing and the Datacenter

4

Locating Things Online

- **Addresses** describe a location
 - A map is OK
 - IP addresses, street addresses
- **Names** are mapped to addresses
 - Domain names; PO Boxes; email addresses
- **Content-based naming** uses the actual content to find the destination
 - Basis of publish/subscribe and peer-to-peer
- What about:
 - Phone #s?
 - Search terms
 - Twitter messages?
 - @username?
 - #topic?

5

Domain Name System

- The Internet's routers know only IP addresses
 - Fine for routers, not great for humans
- DNS translates domain names to IP addresses
 - `google.com` => `74.125.95.99`
 - (or something similar)
 - The "phone book" of the internet
- When DNS goes down, the Web suffers a grievous, not quite deadly, blow


6

DNS Design

- DNS is a globally-accessible database of name/IP pairs. Some requirements:
 - Needs to be up all the time
 - Continuously updated by many parties
 - Must be accurate; errors prevent connections
 - Serves massive query load
 - Needs distributed administration
 - Source of political & commercial disputes
 - Potential terrorism target
- Overall: utterly shocking that it works

7


The Namespace



- Tree structure, 1-63 chars per node
- *Fully-qualified domain name* is leaf-to-root name for path
- Only *DNS root* has no parent

8


The Namespace



- Nodes grouped into administrative *zones*
 - E.g., root, com, google.com
- Each *zone* served by *authority servers* (aka *authoritative name servers*)
- AS can delegate subdomains to other ASes

9


The Namespace



- Servers are primary or secondary
 - Primary ASes are given content by admins
 - Secondary ASes grab from primaries
- A domain registrar inserts your name into the primary AS for .com (or .net, .org, etc)

10

The Namespace



- Purchased domain names inserted into 1 primary, 1 secondary (in case of primary failure)

11

The Data

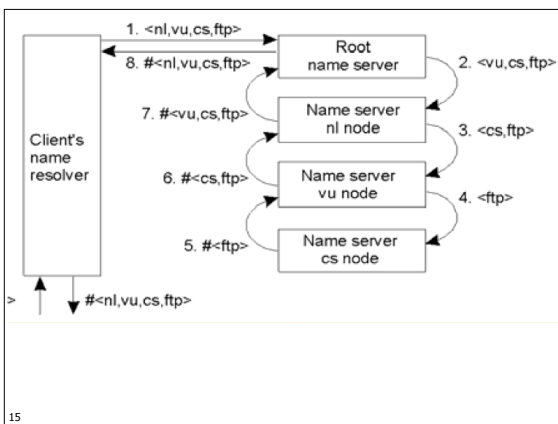
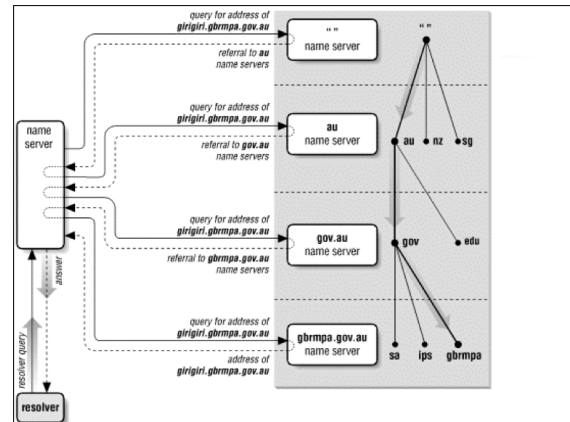
- Content of DNS consists of *resource records*
 - Host/IP and IP/Host are most popular type
 - Also: zone info, email servers, other info
- Most DNS activity is this:
 - DNS client sends request to AS; receives resource record in response
- DNS clients built into network libraries
 - Clients use UDP (not TCP) to grab data
 - If your connection is taking a long time to establish, possibly waiting for DNS

12

Resolution

- Single AS may not be enough for a FQDN

13



Caching

- There may be a chain of caching DNS servers between client and AS
 - There's one for CSE
 - Caches DNS requests; answers new requests from cache whenever possible
- Great! But when data changes?
 - Each RR has time-to-live (TTL) in seconds
 - Counts down from moment AS emits RR
 - Caches and clients must throw out RRs with expired TTLs

16

Root Nameservers

- Responsible for locating the TLD name servers (.com, .net, .org, ...)
- Thirteen root name servers in world
 - Their locations are hard-coded in resolving DNS servers
 - Due to caching, involved in few queries

17

Administration!

- Until 1999, all US TLDs run by IANA
 - (At the time, IANA = Jon Postel)
- Now, non-profit ICANN administers set of for-profit registrars (under contract with gov't)
- For a time, looked like ICANN might be xfer'ed to United Nations. Probably won't be

18

Baffling Corner Cases

- DNS is not very clean; lots of weird holes you don't see in, say, TCP
 - ASes use RRs to describe zone children and parents
 - Parent & child often disagree about relationship, forcing client to do something... reasonable
 - Letter case is supposed to be preserved, but cache implementations mess this up
 - How many requests can a single client message make? In principle, many! In reality, one
- Yet totally world-beating

19


Akamai

- Stores images, videos, at many locations throughout the world
 - Large read-only data stored close to client
 - Reduce latency to faraway datacenter
 - Reduce bandwidth costs by sending data only a short distance
- The 1st CDN (content-delivery network)
 - You might think the image is from Yahoo, but Akamai is serving it
 - Yahoo pays Akamai to do so
- Built on top of giant DNS hack!

20

How it works: step 1


- Akamai places servers across country



21

How it works: step 2


- Akamai gets client data (videos, imgs), copies to all nodes in network



22

How it works: step 3


- How does client find nearby data?
 - Network conditions can change



23

How it works: step 3

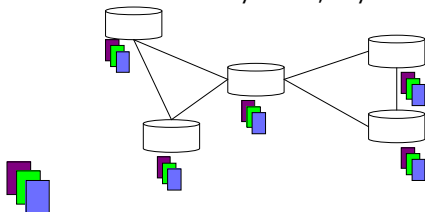
- Client rewrites URLs, uses new URLs
 - `http://yahoo.com/...` => `http://akadns.akamai.net/...`



24

How it works: step 3

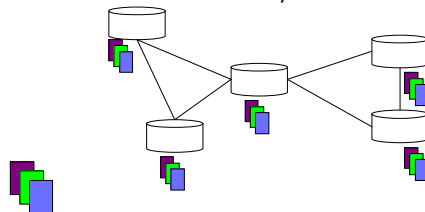
- URL contains Akamai-controlled name
 - Specialized Akamai DNS server resolves name to nearby server; tiny TTL



25

How it works: step 3

- Resolution strongly dependent on
 - Location of client; network conditions; load on Akamai servers; traffic estimation errors



26

Akamai Recap

- Uses custom DNS servers to dynamically direct clients to right caching servers
- How would you implement Akamai without DNS?

27

DNS Challenges

- Shocking scale
 - Every device, service, etc
 - Spam-filtering email requires 10+ DNS lookups per message
- Extreme security vulnerability
 - What if someone can remap `google.com`?
 - What about just denial-of-service?

28

DNS Cache Poisoning

- Bad info is inserted into a DNS server and cached
 - Both inadvertent and malicious
 - How could one attack CSE's DNS server?
- Remedies
 - DNSSEC requires that DNS entries be cryptographically signed
 - If you use HTTPS/SSL/TSL, problem mitigated

29