

Lecture 22

GFS and MapReduce



Administration

- Last call for the datacenter tour!
 - Thurs @2:30PM, or Fri @10:30AM
 - Sign up at the Ctools-posted URL
 - Directions have been posted
 - Take a bus, drive, or get a ride
- Alpha Release this Friday @ 6PM
 - Please keep your sites up till 10AM Sat
- Midterm #2
 - Still being graded
 - Early results: better than MT#1

2

GFS and MapReduce

- Search engines were among first apps to address Web's scale directly
 - We spoke about distributing search queries
 - Also, distributing index construction
- Search engines were first, but not last
 - Social networks
 - Web-analysis and intelligence
 - Log processing
- Traditional data systems inappropriate; Google first to build somewhat general solutions

3

A Little Math

- How hard is indexing the Web really?
 - Say, 10B pages in 1 week
 - 10B in ~10K minutes
 - 1M pages every minute!
 - Disk scan rate 50MB/sec
 - Page takes average 25kb
 - 2000 pages/sec, or,
 - 120,000 pages/min *just to read off disk*
- What if we wanted to do it in a day?
 - An hour?
- Want to use clusters of 100s-1000s of machines, with local disks and some network

4

Distributed Programming

- Unlike distributed databases, many CPUs needed. 1000s, not dozens
 - Programmer cannot know how many machines at program-time or runtime
 - Even so, job is very long-lasting compared to most db queries
 - Machines die, machines depart; job must survive
- Also, cannot express program in SQL

5

MapReduce

- MapReduce system provides:
 - Automatic parallelization & distribution
 - Fault-tolerance
 - Status & monitoring tools
 - Clean abstraction for programmers

6

MapReduce

- Many data programs can be written as *map* and *reduce* functions
- *map* transforms **key, value** inputs into new **key', value'**
 - **Map(k, v) => (k', v') list**
- *reduce* receives all the vals for a given key' and can output to disk file
 - **Reduce(k', v' list) => (out-key, out-val) list**

7

Example: Filter

- **Map(k, v) =>**
 - **if (prime(k)) then emit(k, v)**
- ("foo", 7) => ("foo", 7)
- ("bar", 10) => *nothing*

8

Example: Sum

- **Reduce(k, vals) =>**
 - **sum = 0**
 - **foreach int v in vals:**
 - **sum += v**
 - **emit(k, sum)**
- ("A", [42, 100, 312]) => ("A", 454)
- ("B", [12, 6, -2]) => ("B", 16)

9

Word counting

- **Map(lineNo, textStr)**
 - **for each word w in textStr:**
 - **emit(w, 1)**
- **Reduce(outKey, interm-vals)**
 - **int result = 0**
 - **for each v in interm-vals:**
 - **result += v**
- **Emit(output-key, result)**

10

Applications

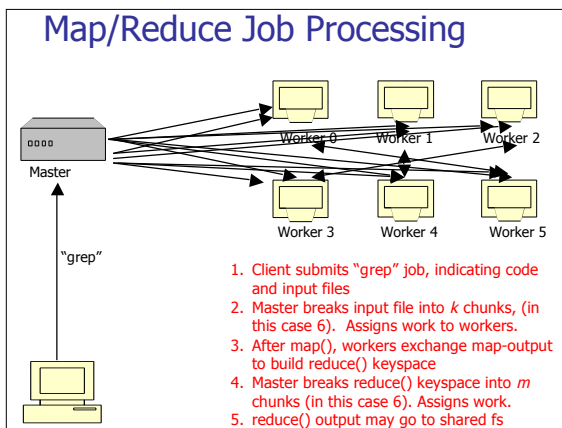

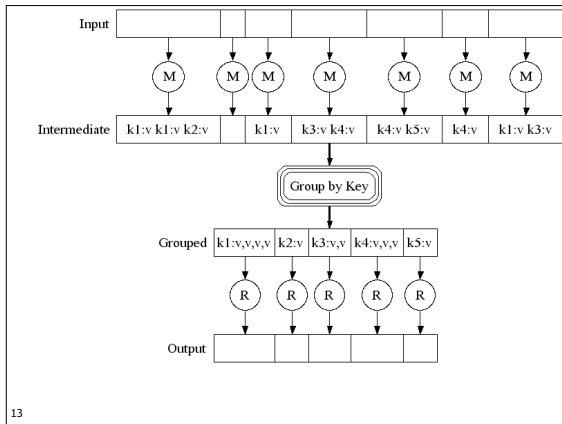
- What else can be an MR program?
 - URL counting in logs, Reverse Web graph
 - Inverted index construction, Sorting
 - Massive image conversion, others

11



Shuffle/Sort


- What happens between map & reduce?
 - Data collated and grouped for map
- Execution goes as follows:
 - Break input into M chunks
 - Process each chunk w/ *map* process
 - Group map outputs into R chunks
 - Process each chunk w/ *reduce* process
 - *reduce* fn's outputs go to disk

12



- ### A few nice features
- What about slow, not dead, machines?
 - Speculative execution for stragglers
 - What about data placement?
 - Use distributed filesystem across cluster disks; start tasks where the target data already lies
 - Isn't the intermediate data size large?
 - Use a "local reducer" called a Combiner at each map
 - Compress data between map and reduce
- 16






- ### Performance
- The TeraSort benchmark measures time to sort 1TB (10B 100-byte records)
 - In 1998, record holder did it in 151 mins
 - In 2009, Yahoo/Hadoop MapReduce did it in 209 seconds
 - Then a few months later, Google did it in 68 seconds (on about 1,000 machines)
 - MapReduce also holds the record for sorting a petabyte
 - (That's 10 trillion 100-byte records)
 - 6 hours, 2 mins on 4,000 machines
- 20

- ### Intermission
- MapReduce has become extremely trendy. Any problem, no matter how unrelated, is now treated as a good target for MapReduce programming
 - That said, it has its uses. For a different opinion, do a search for "MapReduce: A Major Step Backwards"
- 21

- ### Google File System
- The distributed filesystem that stores data across a MapReduce cluster
 - As with MR, Google had strange requirements and came up with strange design
 - Has not proved as generally-useful as MapReduce
- 22

- ### Weird Requirements
- High component failure rates
 - Inexpensive commodity components fail all the time
 - "Modest" number of HUGE files
 - Just a few million
 - Each is 100MB or larger; multi-GB files typical
 - Files are append-only
 - Large streaming reads
 - High sustained throughput favored over low latency
- 23

- ### Design solution
- Files stored as chunks
 - Much larger size than most filesystems (default is 64MB)
 - Reliability through replication
 - Each chunk replicated across 3+ *Chunkservers*
 - Single master coordinates access, metadata
 - Simple centralized management
 - No data caching
 - Little benefit due to large data sets, streaming reads
 - Familiar interface, but customize the API
 - Simplify the problem; focus on distributed apps
- 24



Design Details

- Chunks can be copied, replicated
- Chunkservers hold and serve chunks
- Master holds metainfo
 - Filename → chunk list
 - Chunk → chunkserver-location
- Chunkservers report in to master every few seconds
- Data/task co-location is extremely important

25

